

数据库week11

15.2

Consider the bank database of Figure 15.14, where the primary keys are underlined, and the following SQL query:

```
branch(branch_name, branch_city, assets)
customer(customer_name, customer_street, customer_city)
loan(loan_number, branch_name, amount)
borrower(customer_name, loan_number)
account(account_number, branch_name, balance)
depositor(customer_name, account_number)
```

Figure 15.14 Bank database.

```
select T.branch_name
from branch T, branch S
where T.assets > S.assets and S.branch_city = "Brooklyn"
```

Write an efficient relational-algebra expression that is equivalent to this query. Justify your choice.

Expression:

$$\Pi_{T.branch_name}((\Pi_{branch_name, assets}(\rho_T(branch))) \bowtie_{T.assets > S.assets} (\Pi_{assets}(\sigma_{branch_city='Brooklyn'}(\rho_S(branch)))))$$

By eliminating the unneeded attributes from both the operands, we can decrease the blocks we will transfer.

15.3

Let relations $r_1(A, B, C)$ and $r_2(C, D, E)$ have the following properties: r_1 has 20,000 tuples, r_2 has 45,000 tuples, 25 tuples of r_1 fit on one block, and 30 tuples of r_2 fit on one block. Estimate the number of block transfers and seeks required using each of the following join strategies for $r_1 \bowtie r_2$:

$$b_{r_1} = \frac{20000}{25} = 800$$

$$b_{r_2} = \frac{45000}{30} = 1500$$

a. Nested-loop join.

If r_1 is the outer relation, we need $n_{r_1} * b_{r_2} + b_{r_1} = 30,000,800$ block transfers and $n_{r_1} + b_{r_1} = 20800$ disk seeks.

If r_2 is the outer relation, we need $n_{r_2} * b_{r_1} + b_{r_2} = 36001500$ block transfers and $n_{r_2} + b_{r_2} = 46,500$ disk seeks.

b. Block nested-loop join.

If r_1 is the outer relation, we need $\lceil \frac{b_{r_1}}{M-2} \rceil * b_{r_2} + b_{r_1}$ block transfers and $2 * \lceil \frac{b_{r_1}}{M-2} \rceil$ disk seeks.

If r_2 is the outer relation, we need $\lceil \frac{b_{r_2}}{M-2} \rceil * b_{r_1} + b_{r_2}$ block transfers and $2 * \lceil \frac{b_{r_2}}{M-2} \rceil$ disk seeks.

c. Merge join.

- Assuming that r_1 and r_2 are not initially sorted on the join key and $b_b = 1$, the total sorting transfers cost inclusive of the output T_Sort is

$$T_Sort(r_i) = b_{r_i}(2\lceil \log_{M-1}(b_{r_i}/M) \rceil + 1)$$

the total sorting transfers cost inclusive of the output S_Sort is

$$S_Sort(r_i) = 2\lceil \frac{b_{r_i}}{M} \rceil + b_{r_i}(2\lceil \log_{M-1}(b_{r_i}/M) \rceil - 1)$$

So we need $T_Sort(r_1) + T_Sort(r_2) + b_{r_1} + b_{r_2}$ block transfers and $S_Sort(r_1) + S_Sort(r_2) + b_{r_1} + b_{r_2}$ disk seeks.

- If r_1 and r_2 are sorted on the join key and $b_b = 1$

So we need $b_{r_1} + b_{r_2}$ block transfers and $b_{r_1} + b_{r_2}$ disk seeks.

d. Hash join.

Assuming that the $b_b = 1$.

- If r_1 as probe input, r_2 as build input

If doesn't need recursive partitioning, we need $3(b_{r_1} + b_{r_2}) = 6900$ disk accesses and $2(b_{r_1} + b_{r_2}) = 4600$ seeks, else, we need $2(b_{r_1} + b_{r_2}) \lceil \log_{M-1} \frac{b_{r_2}}{M} \rceil + b_{r_1} + b_{r_2}$ disk accesses and $2(b_{r_1} + b_{r_2}) \lceil \log_{M-1} \frac{b_{r_2}}{M} \rceil$ seeks.

- If r_2 as probe input, r_1 as build input

If doesn't need recursive partitioning, we need $3(b_{r_1} + b_{r_2}) = 6900$ disk accesses and $2(b_{r_1} + b_{r_2}) = 4600$ seeks, else, we need $2(b_{r_1} + b_{r_2}) \lceil \log_{M-1} \frac{b_{r_1}}{M} \rceil + b_{r_1} + b_{r_2}$ disk accesses and $2(b_{r_1} + b_{r_2}) \lceil \log_{M-1} \frac{b_{r_1}}{M} \rceil$ seeks.